



Model-Based Reinforcement Learning

Michael L. Littman

Rutgers University

Department of Computer Science

Rutgers Laboratory for Real-Life Reinforcement Learning




Topics

- Introduction
 - MDPs
 - Reinforcement Learning
- Model-based RL
- Efficient Exploration
 - PAC-MDP
 - KWIK
- Bayesian RL
 - Near-Bayesian
 - PAC-MDP
- Planning
 - Nesting approaches
 - UCT
- Model-based People?

Start With Game...

- up
- down
- left
- right
- A
- B



Find The Ball: Elements of RL

In reinforcement learning:

- agent interacts with its environment
- perceptions (state), actions, rewards [repeat]
- task is to choose actions to maximize rewards
- complete background knowledge unavailable

Learn:

- which way to turn
- to minimize time
- to see goal (ball)
- from camera input
- given experience.





Problem To Solve


Three core issues in the dream RL system.

- generalize experience
 - use knowledge gained in similar situations
 - "learning"
- sequential decisions
 - deal properly with delayed gratification
 - "planning"
- exploration/exploitation
 - must strike a balance
 - unique to RL?

Markov Decision Processes

Model of sequential environments (Bellman 57)

- n states, k actions, discount $0 \leq \gamma \leq 1$
- step t , agent informed state is s_t , chooses a_t
- receives payoff r_t ; expected value is $R(s_t, a_t)$
- probability that next state is s' is $T(s_t, a_t, s')$




$$Q(s,a) = R(s,a) + \gamma \sum_{s'} T(s,a,s') \max_{a'} Q(s',a')$$

- Optimal behavior is $a_t = \operatorname{argmax}_a Q(s_t, a)$
- R, T unknown; some experimentation needed


Find the Ball: MDP Version

- Actions: rotate left/right
- States: orientation
- Reward: +1 for facing ball,
0 otherwise



Families of RL Approaches

policy search value-function based model based



More direct use,
less direct learning

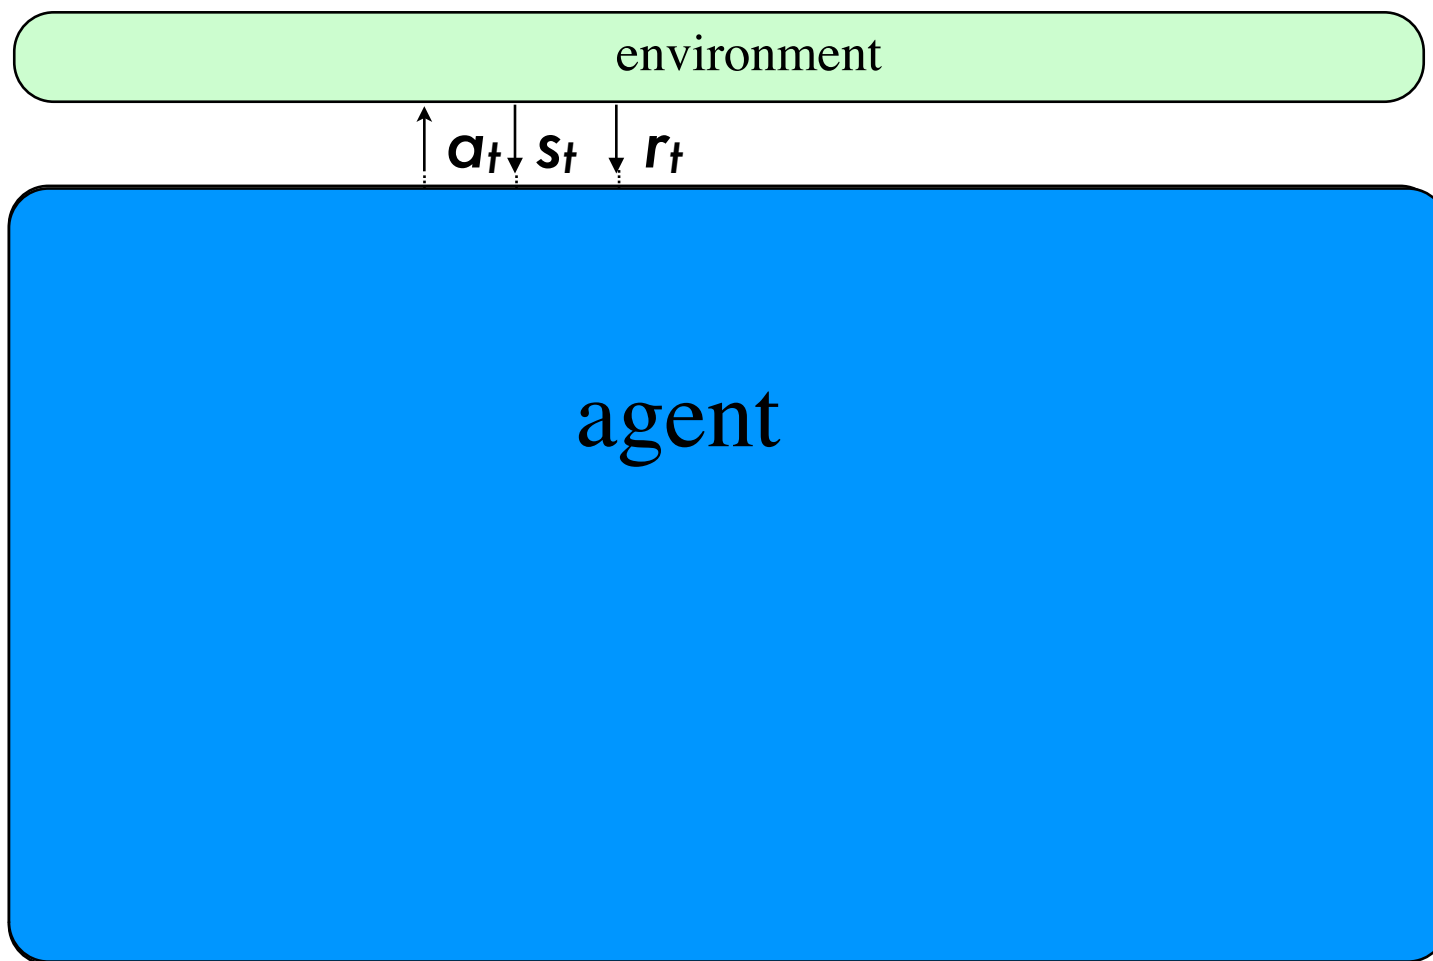
Search for
action that
maximizes
value

Solve Bellman
equations

More direct learning,
less direct use



Model-based RL Schematic



PAC-MDP Reinforcement Learning




PAC: Probably approximately correct (**Valiant 84**)

Extended to RL (**Fiechter 95, Kakade 03, etc.**).

- Given $\epsilon > 0$, $\delta > 0$, k actions, n states, Υ .
- We say a strategy makes a mistake each timestep t s.t. $Q(s_t, a_t) < \max_a Q(s_t, a) - \epsilon$.
- Let m be a bound on the number of mistakes that holds with probability $1 - \delta$.
- Want m poly in $k, n, 1/\epsilon, 1/\delta, 1/(1 - \Upsilon)$.

Must balance exploration and exploitation!

Model-based Can Be PAC-MDP



- Behavior differs depending on assumption

	truth: ? = low	truth: ? = high
assume: ? = low	ignore ?, optimal	ignore ? suboptimal!
assume: ? = high	visit ?, explore	visit ?, optimal

← No PAC-MDP
guarantee

← PAC-MDP if not too
much exploration



Model-driven Exploration

- A *model* is $\mathbb{R}(s,a)$ and $\mathbb{T}(s,a,s')$.
- Model-based approach:
 - learn a model of the environment (approximately, distinguishing known/unknown transitions).
 - augment model w/ bonus for unknown transitions.
 - plan behavior wrt the augmented model.
 - repeat

Key that learner “knows what it knows” (KWIK).

What learning setting is appropriate?

3 Models for Learning Models

- PAC: Inputs drawn from a fixed distribution. Observe inputs. For future inputs from the distribution,

Not PAC-MDP. iid assumption implies that learner cannot improve (change) behavior!

- Mistake bound: Inputs presented online. For each, predict output. If mistake, observe label. If more than m mistakes,

Not PAC-MDP. Mistakes mean that a high reward can be assumed low—suboptimal.

- KWIK: Inputs presented online. For each, can predict output or say "I don't know" and request label. No mistakes, but can say "I don't know" m times.

Can be PAC-MDP...

no mistakes



KWIK-Rmax Proof

- Provides PAC-MDP guarantee in flat MDPs ([Kearns & Singh 02](#), [Brafman & Tennenholtz 02](#)).
 - Key ideas:
 - Simulation lemma: Optimal actions for approximate model near-optimal in real model.
 - Explore or exploit lemma: If can't reach unknown states quickly, can achieve near-optimal reward.
 - Unflat: factored dynamics ([Kearns & Koller 99](#)), metric spaces ([Kakade et al. 03](#)), KWIK ([Li 09](#)).
- Time to learn depends on KWIK bound.*



KWIK Learn a Probability

- Given m trials, x successes, $\hat{p} = x/m$
- Hoeffding bound:
 - Probability of an empirical estimate of a random variable in the range $[a,b]$ based on m samples being more than ϵ away from the true value is bounded by $\exp\left(-\frac{2m\epsilon^2}{(b-a)^2}\right)$
- So, can KWIK learn a transition probability:
 - say “I don’t know” until m is big enough so that \hat{p} is ϵ -accurate with probability $1-\delta$.

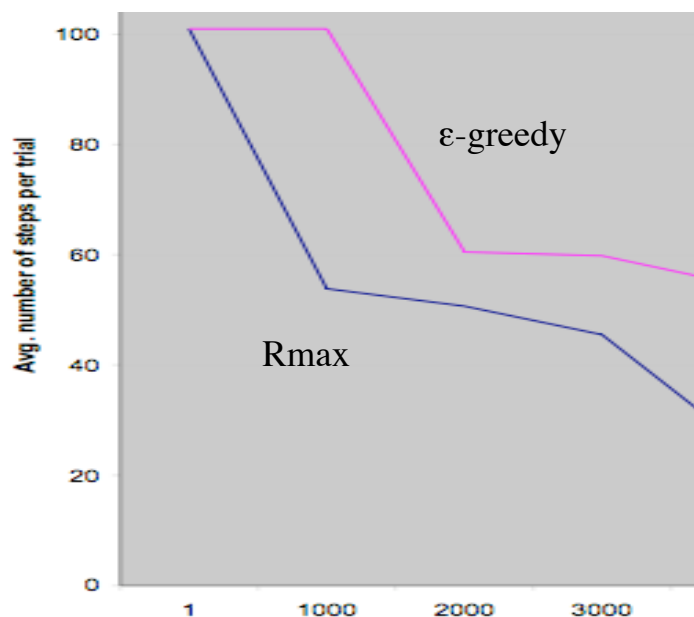


Other Things to KWIK Learn

- coin probability
- vector of outputs, each KWIK learnable
 - multinomial probability (dice learning)
- mapping from input partition to outputs, partitions known, mappings KWIK learnable
 - That's a standard transition function (s, a to vector of coins) (Li, Littman, Walsh 08).
- Also, union of two KWIK learnable classes.

R_{MAX} Speeds Learning

Task: Exit room using bird's-eye state representation.





Details: Discretized 15x15 grid x 18 orientation (4050 states);
6 actions: forward, backward, turn L, turn R, slide L, slide R.

(Nouri)

Generalizing Transitions

- Flat MDPs, states viewed as independent.
 - Transition knowledge doesn't transfer.
- DBN representation shares structure.
 - Learn components independently, KWIK!



Less to learn, faster to behave well.


Continuous-state DBN



(Nouri)

World of Objects

- Objects in taxi:
 - taxi (location)
 - passenger (location/in taxi)
 - walls (location)
 - destination (location)
- Not *states* or state *features*, instead try *objects* and object *attributes*.
- Model: What happens when objects interact?
- More “human like” exploration.



Comparing Taxi Results

- North, not touchN(taxi, wall) \rightarrow taxi.y++
- Drop, pass.in, touch(taxi, dest) \rightarrow \neg pass.in
- KWIK bound: poly in types (exp in condition)
- Taxi: How long until optimal behavior?

Exploration style	Algorithm	# of steps
ϵ greedy	Q-learning	47157
count on states	Flat Rmax	4151
count on features	Factored Rmax	1839
count on interaction	Objects	143
whatever people do	People	50

Pitfall!



A childhood dream fulfilled... (Diuk, Cohen)



Structure Learning in DBNs

- Unknown structure fundamentally different.
- How can you keep statistics if you don't know what they depend on?
- Can be solved using a technique for a simpler "hidden bit" problem:
 - n -bit input, one bit (unknown) controls output
 - one output distribution if bit is on, another if off
 - Find DBN structure by same idea: one parent set controls output...



Hidden-Bit Problem

Assume the simpler deterministic setting.

Output is copy or flip of one input.

- 0110 → 0 1101 → 1 0000 → 1
- 1101 → 1 0011 → 0 1111 → 0
- 1000 → 1 1110 → 0 1100 → 1

Is it 0, 1, or “I don’t know”?

If noisy, can’t predict with each bit position separately, don’t know which to trust. Can learn about all 2^n bit patterns separately, but that’s too much.



Hidden-bit Problem via KWIK

- Can observe predictions to figure out which of k “adaptive meteorologists” to trust (Strehl, Diuk, Littman 07; Diuk et al. 09).
- Solvable with bound of $O\left(\frac{k}{\epsilon^2} \ln \frac{k}{\delta}\right) + \sum_{i=1}^k \zeta_i\left(\frac{\epsilon}{8}, \frac{\delta}{k+1}\right)$
- By considering all k -size parent sets, get a structure-learning algorithm with a KWIK bound of

$$\kappa = O\left(\frac{n^{D+3}AD}{\epsilon^3(1-\gamma)^6} \ln \frac{nA}{\delta} \ln \frac{1}{\epsilon(1-\gamma)}\right)$$

Artificial Stock Example

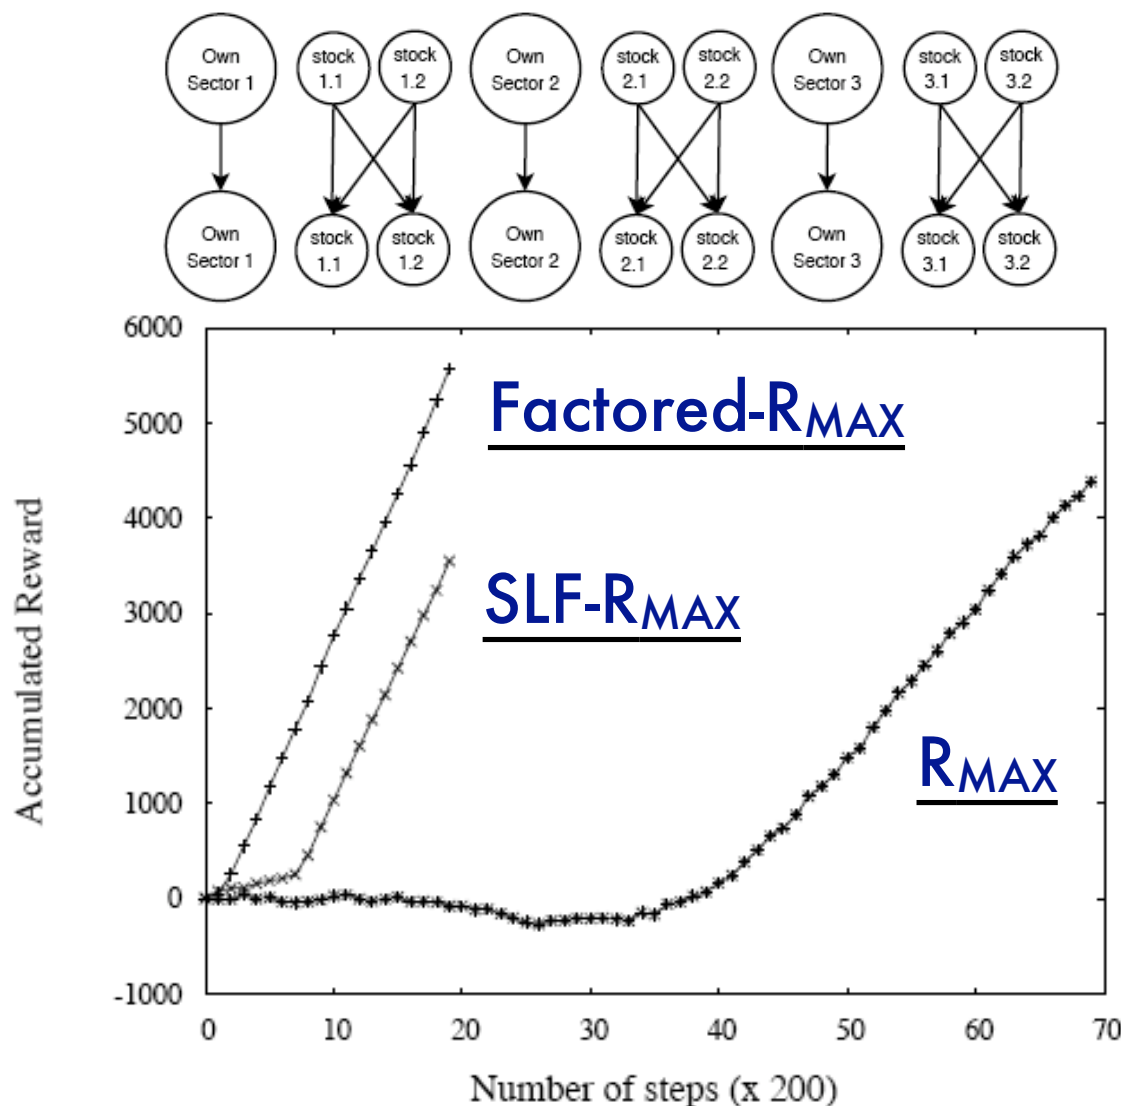
Discovers the structure and exploits it much faster than R_{MAX} can learn the MDP.

Factored- R_{MAX} :

Knows DBNs

SLF- R_{MAX} : Knows size of parent sets

R_{MAX} : It's an MDP






Many Learnable Problems

Many hypothesis classes KWIK learnable:

- coin flip probability
- Dynamic Bayes net probabilities given graph
- k Dynamic Bayes net
- k Meteorologist problem
- k -CNF
- k -depth decision tree
- unions of KWIK-learnable classes
- k feature linear function

Beyond Worst Case

- KWIK learns specific hypothesis class.
 - If too broad, learning too slow.
 - If too narrow, learning fails.
- Bayesian perspective:
 - Start with a prior over models.
 - Maintain a posterior.
 - Optimize for *probable* instead of just *possible*.
 - Similar states have similar dynamics. Probably.
- Bayesian view can drive exploration.





Bayes Optimal Exploration

- With a Bayesian representation of models, we can plan in the space of *posteriors*.
 - Can use posterior to evaluate the likelihood of any possible outcome of an action.
 - Can model how that outcome changes posterior.
 - Can choose actions that truly maximize expected reward: No artificial distinction between exploring and exploiting or learning and acting!
- Hideously intractable except in special cases (bandits, short horizons).

Concrete Example


- MDP has one state, 3 actions (bandit)

- $X: \{.7 \ .1 \ .8\}$, $Y: \{.8 \ .6 \ .7\}$, $\gamma = 0.8$

- Prior: $\langle .50, .50 \rangle$ ($1/2 \ X, 1/2 \ Y$)

$\langle .50, .50 \rangle$

$(+.750)$



Concrete Example


- MDP has one state, 3 actions (bandit)

- $X: \{.7 \ .1 \ .8\}$, $Y: \{.8 \ .6 \ .7\}$, $\gamma = 0.8$

- Prior: $\langle .50, .50 \rangle$ ($1/2 \ X, 1/2 \ Y$)


$\langle .50, .50 \rangle$

$(+.750)$



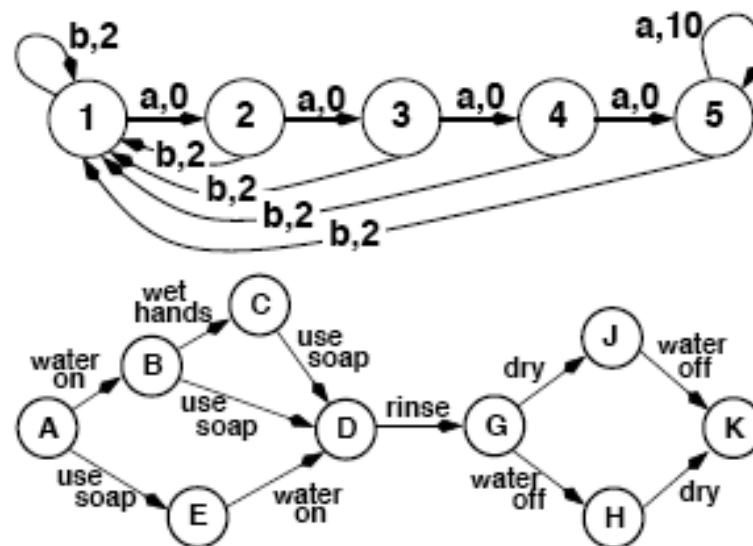
Representing Posteriors

- $T: s, a \rightarrow$ multinomial over states
- If independent for each s, a :
Dirichlet!
- Keep counts for each observed outcome.
- Can recover uncertainty in overall estimate.
- Unlike example, distribution over an infinite set.



Bayes Optimal Plans

- Many attempts (Duff & Barto 97; Dearden et al. 99)
- State of the art, BEETLE (Poupart et al. 06).
 - Latest ideas from solving continuous POMDPs
 - α functions are multivariate polynomials + PBVI
 - Can exploit “parameter tying” prior.
 - Near optimal plan in “combination lock”.
 - Less optimal in bigger problem.
 - Planner outputs exploration scheme.






Near Bayes Optimal Behavior

- Recall PAC-MDP, whp makes few mistakes.
- Near Bayesian: mistakes are actions taken with values far from Bayes optimal.
- Bayesian Exploration Bonus (Kolter & Ng 09) keeps mean of posterior and adds $1/n$ bonus to actions taken n times.
 - BEB is computationally simple.
 - BEB is Near Bayesian.
 - BEB is not PAC-MDP, though...




Bayes Optimal Not PAC-MDP

- Examples where Bayes optimal does not find near optimal actions (**Kolter & Ng 09; Li 09**)



- Bayes optimal approach chooses a_1 forever even though not ϵ -optimal (for $p > 2\epsilon$).
- So, even if the two models are equally likely, Bayes optimal doesn't bother learning about the better option!


Inherent Conflict...



- PAC-MDP: Future self gets near-optimal reward.
- Near Bayesian: Current self gets near-optimal reward.
- Human behavior in between? (Hyperbolic discounting.)

PAC-MDP with Bayesian Priors

- With a prior that all similar colored squares are the same, we can bound the chance generalization will lead to sub-optimality.
- Idea: Don't worry about it if it's small!



$X: \{.7 \ .1 \ .8\}, Y: \{.8 \ .6 \ .7\}$

$\epsilon=0.0001, \delta=0.05$

$\langle .99, .01 \rangle$

R is near optimal whp




BOSS: Algorithmic Approach

- Optimism under uncertainty, not Bayes optimal
 - Sample models from the posterior.
 - Stitch together into a meta-MDP.
 - Solve to find optimal behavior: best of sampled set
 - Act accordingly until something new learned.
- If set big, near optimality whp ([Asmuth et al. 09](#))
- Several ideas appear to be viable here

$$O\left(\frac{SAB}{\epsilon(1-\gamma)^2} \ln \frac{1}{\delta} \ln \frac{1}{\epsilon(1-\gamma)}\right)$$

BOSS in Structured Maze

- To learn in maze:
 - Chinese Restaurant Process prior
 - Finds (empirical) clusters
 - Outperforms Rmax, 1-cluster RAM-Rmax



- Fewer than states
- Fewer than types
- Some types grouped
- Rare states nonsense




Computation Matters


- Learning/exploration can be made efficient
 - model-based RL
 - PAC-MDP for studying efficient learning
 - KWIK for acquiring transition model
- Planning “just” a computational problem.
 - But, with powerful generalization, can quickly learn accurate yet intractable models!
 - Something needs to be done or the models are useless.

"Nesting" RL Approaches


policy search




value function




model-based



policy search inside model-based





value function inside model-based




Function Approximation

A natural NIPS/model-based RL connection.
Use your favorite regression algorithm for T .
Use T as a simulator and run your favorite RL.
(Moore, Atkeson & Schaal 95; Jong & Stone 06)



Example: Autonomous Flight

- Outer approach: Model-based RL.
 - Experts parameterize model space.
 - Parameters learned quickly from expert demonstration (no exploration needed).
- Resulting model very high dimensional (S,A)
- Inner approach: Policy-search RL.
 - Experts parameterize space of policies.
 - Offline search finds excellent policy on model.
 - Methodology robust to error in model.
- Learns amazing stunts (Ng et al. 03).






Tricks and Treats



Stanford University Autonomous Helicopter


Fitted Value Iteration

- Represent value function via anchor points and local smoothing (Gordon 95)
- Some guarantees if points densely sampled (Chow & Tsitsiklis 91)
- Combined with KWIK learning of model (Brunskill et al. 08)




UCT: Upper Conf. in Trees

- Narrow, deep game-tree search via bandits
(Kocsis & Szepesvári 06)
- Huge win in Go (Gelly & Wang 06; Gelly & Silver 07)
- Good fit w/ learned model.
 - Just needs to be able to simulate transitions.
 - KWIK-like methods are also “query” based.
- Not much work using it in RL setting.




Linear Models

- Linear value function approaches: LSTD/LSPI (Boyan 99; Lagoudakis & Parr 03, Parr et al. 08).
- Method 1: Learn linear dynamics model from sample. Solve to get (linear) value function.
- Method 2: Learn value function directly.




Do Kids Explore Models?

- Statistics of play sensitive to confounding
- Show kid 2-lever toy (Schulz/Bonawitz 07).
 - Demonstrate levers separately. Kid more interested in new toy.
 - Demonstrate them together. Kid stays interested in old toy.
- Experiment design intractable. KWIK-like heuristic?



Old Toy



New Toy

Do People Explore Models?



xkcd





Wrap-Up

- Introduction
 - MDPs
 - Reinforcement Learning
- Model-based RL
- Efficient Exploration
 - PAC-MDP
 - KWIK
- Bayesian RL
 - Near-Bayesian
 - PAC-MDP
- Planning
 - Nesting approaches
 - UCT
- Model-based People?